

CS420: Operating Systems

Threads

James Moscola
Department of Physical Sciences
York College of Pennsylvania



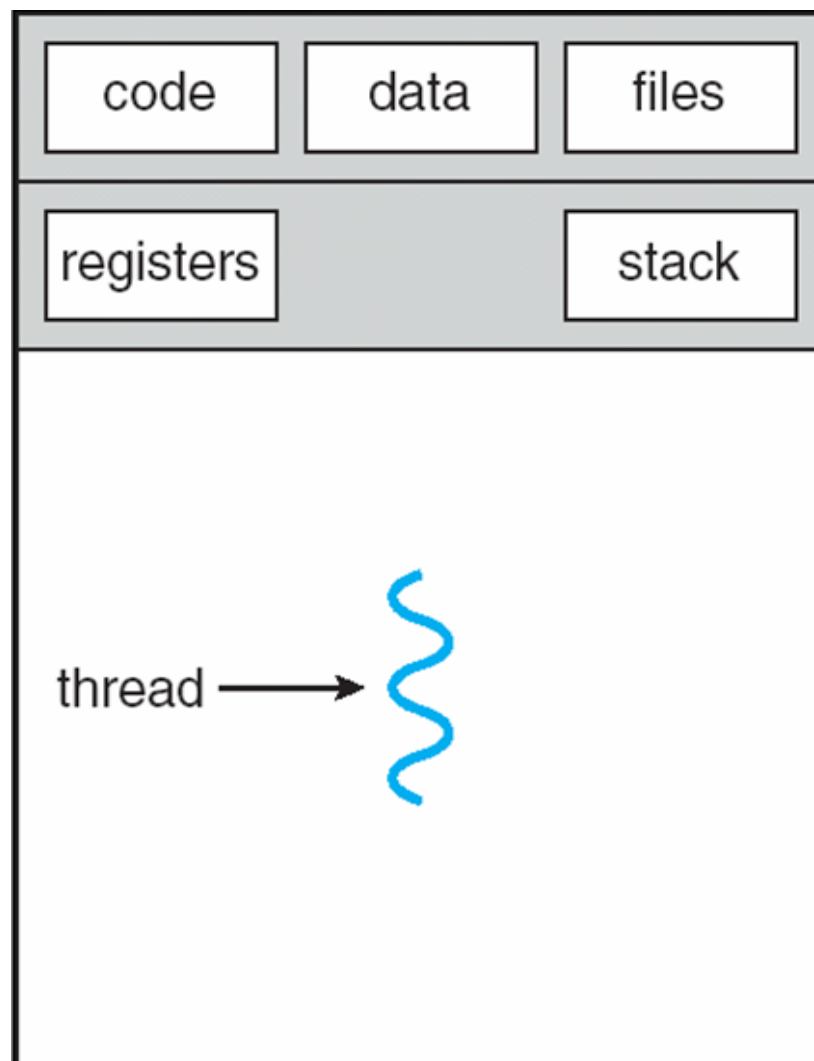
Threads

- **A thread is a basic unit of processing**
 - Has the following components:
 - Thread ID
 - Program counter
 - Register set
 - Stack
 - Shares some resources with other threads in same process
 - Code section
 - Data section
 - OS Resources (e.g. open files, signals)
- **Scheduled by the operating system**

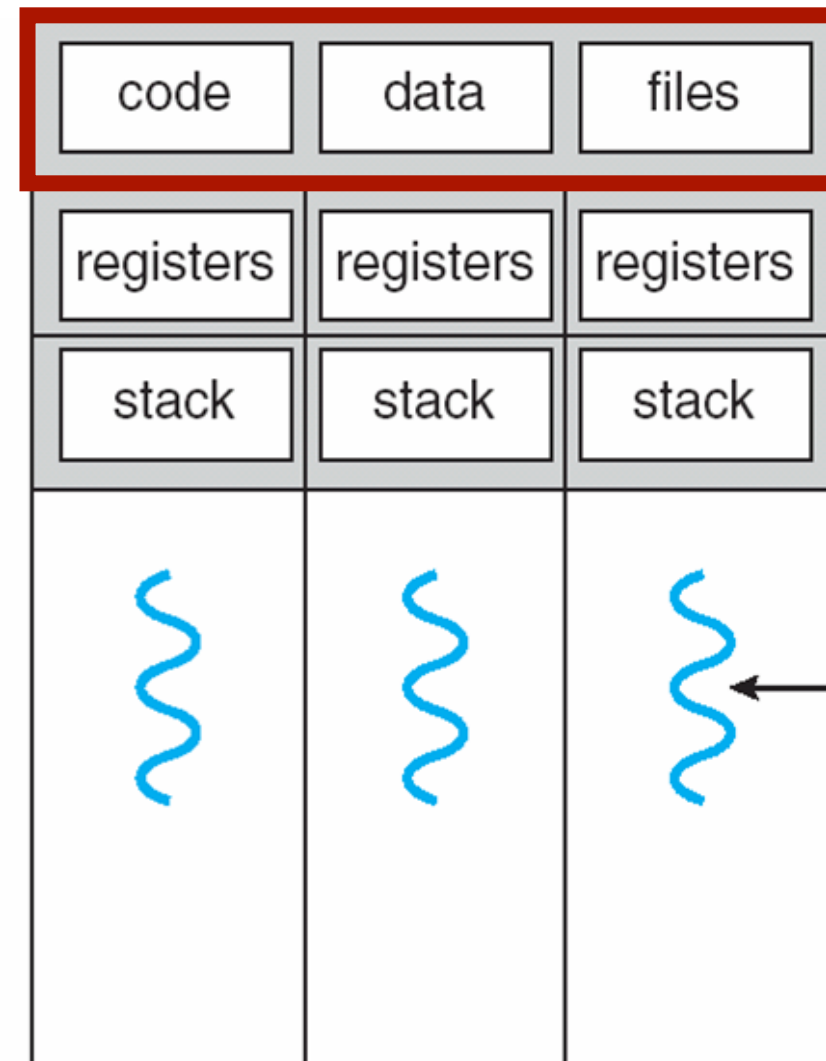
Threads

- **A heavyweight process is a process that has a single thread of control**
 - Can only perform a single task at a time
- **A multi-threaded process is a process that has multiple threads of control**
 - Can perform more than one task at a time
 - Render images
 - Fetch data
 - Update display
 - Check spelling

Single and Multithreaded Processes



single-threaded process



Shared

multithreaded process

Thread vs Process

- **Processes -**

- Independent units of execution
- Each contains its own state information
- Each contains its own address space
- Interact with each other through various IPC mechanisms

- **Threads (within the same process) -**

- Share the same state
- Share the same memory space
- Share the same variables
- Can communicate directly through shared variables
- Share signal handling

Benefits of Multithreaded Programming

- **Responsiveness**

- Interactive applications are more responsive when multithreaded (e.g. a thread for the GUI, another for socket, a third for rendering, etc.)

- **Resource Sharing**

- Unlike processes, threads share memory and resources

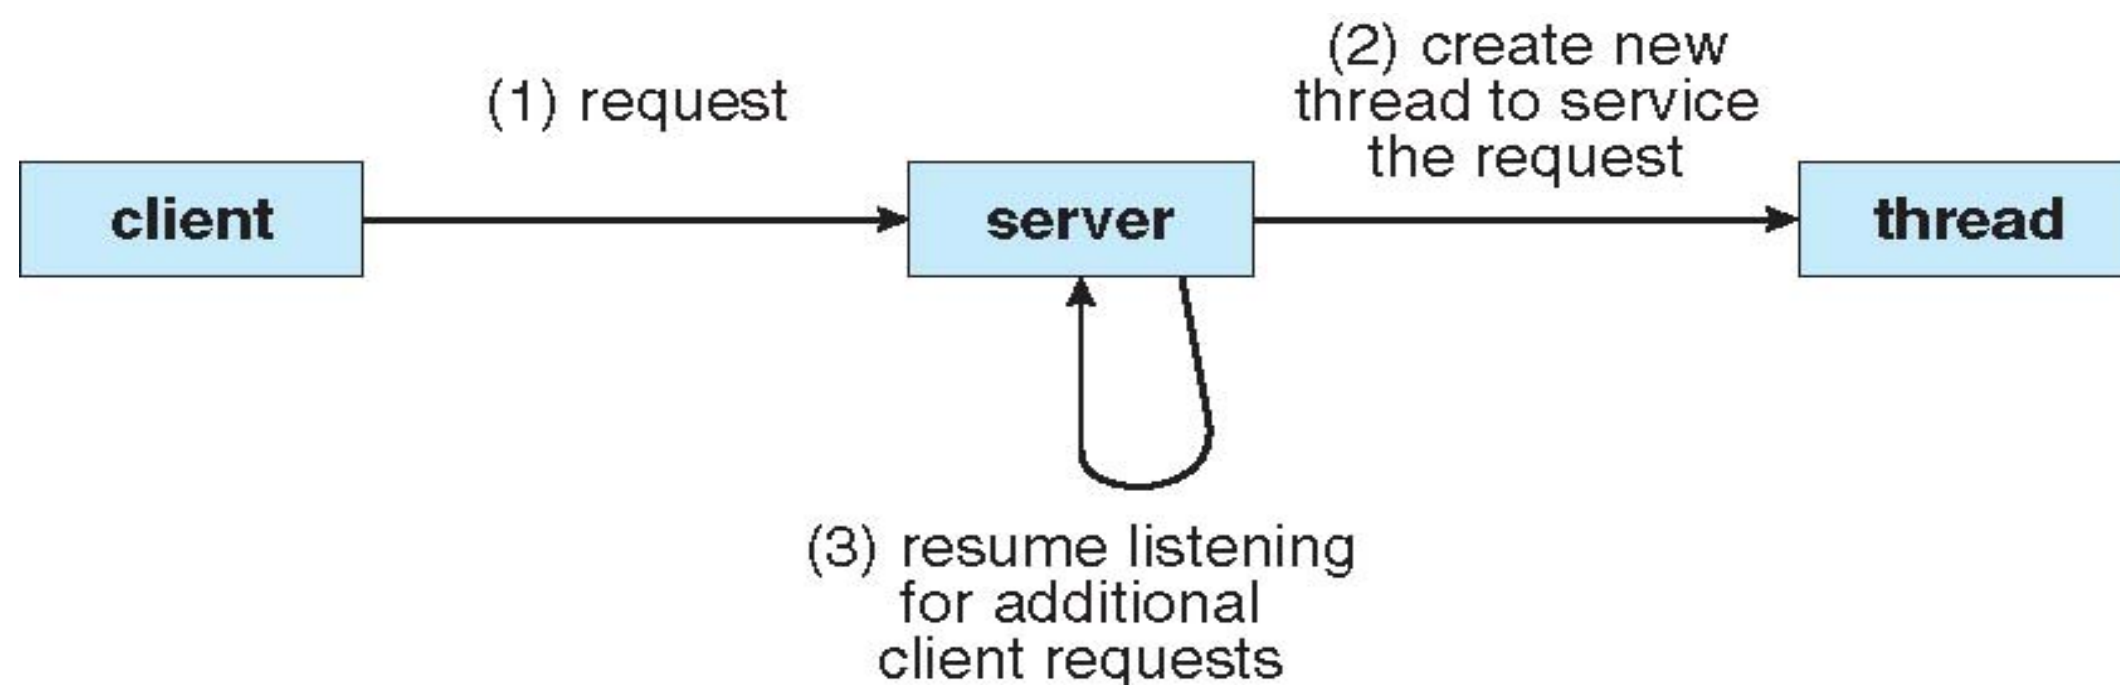
- **Economy**

- Since threads share resources, creating threads and switching between them is more efficient than processes

- **Scalability**

- Multithreading allows for increased parallelism on multicore systems as each thread can run on a different CPU core


Multithreaded Server Architecture



Multicore Programming


- **The performance of a process can be increased by properly threading the process to take advantage of modern multicore CPUs**
- **Multicore systems have challenges not faced in a single-core/single-threaded environment**
 - **Dividing activities** - How can an application be divided into separate, concurrent tasks?
 - **Balance** - How can those tasks be divided in such a way that each does an equal amount of work?
 - **Data splitting** - Can the data for those task be divided for processing on separate CPU cores?
 - **Data dependency** - Are there data dependencies between different task?
 - **Testing and debugging** - What is the best way to debug a multithreaded program with many different execution paths?

Concurrent Execution on a Single-core System



- Only a single thread can execute at a time
- Threads are interleaved so each gets time on the processor

Parallel Execution on a Multicore System



- With multiple cores, threads can be divided over the cores and run in parallel
- May still interleave threads if not enough cores are available for all of the threads

Thread Support


- **Threads may be supported at different levels of the OS**
 - **User threads**
 - Supported above the kernel
 - Managed without kernel support
 - Three main user thread libraries currently in use -
(1) POSIX PThreads (2) Win32 threads (3) Java threads
 - **Kernel threads**
 - Supported by the kernel/operating system
 - Managed by the kernel/operating system
 - Most modern operating systems support kernel threads
(e.g. Windows XP/2000, Solaris, Linux, Tru64 UNIX, Mac OS X)

Multithreading Models

- **There must be a relationship between user level threads and kernel threads**
- **Different models of threading exist to define this relationship**
 - Many-to-One
 - One-to-One
 - Many-to-Many


Many-to-One

- **Many user-level threads mapped to single kernel thread**
 - Examples: Solaris Green Threads, GNU Portable Threads
- **Thread management is done in user space**
- **Entire process will block if any single thread blocks (no other threads will run)**
- **Unable to run multiple threads in parallel on a multiprocessor system**




One-to-One

- **Each user-level thread maps to kernel thread**
 - Examples: Windows NT/XP/2000, Linux, Solaris 9 and later
- **Allows more concurrency**
 - A thread can run when another thread has made a blocking system call
 - Multiple threads can run in parallel on multiprocessor systems
- **Downside, for each thread created, a corresponding kernel thread must also be created**



Many-to-Many Model

- **Allows many user level threads to be mapped to many kernel threads**
 - Avoids blocking of threads when a single thread makes a blocking system call
- **Allows the operating system to create a sufficient number of kernel threads**
- **Reduces the overhead associated with too many kernel threads as was present in the one-to-one model**




Two-level Model

- **Similar to the Many-to-Many model except that it allows a user thread to be bound to a specific kernel thread**

- **Examples include**

- IRIX
- HP-UX
- Tru64 UNIX
- Solaris 8 and earlier



Thread Libraries

- **A thread library provides programmer with API for creating and managing threads**
- **Two primary ways of implementing**
 - Library entirely in user space
 - Kernel-level library supported by the OS
- **Three main user thread libraries currently in use**
 - (1) POSIX PThreads - user-level or kernel-level threads for POSIX-compliant systems
 - (2) Win32 threads - kernel-level threads for Windows systems
 - (3) Java threads

POSIX Pthreads

- **May be provided either as user-level or kernel-level**
- **A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization**
- **API specifies behavior of the thread library, implementation is up to development of the library**
- **Common in UNIX operating systems (Solaris, Linux, Mac OS X)**

Java Threads

- **Java threads are managed by the JVM**
- **Typically implemented using the threads model provided by underlying OS**
- **Java threads may be created in two different ways:**
 - Extending the Thread class
 - Implementing the Runnable interface

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

Pthreads Example (Cont.)

```
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid,&attr,runner,argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Figure 4.9 Multithreaded C program using the Pthreads API.

Win32 API Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */
/* the thread runs in this separate function */

DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;
    /* perform some basic error checking */
    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```

Win32 API Multithreaded C Program (Cont.)

```
// create the thread
ThreadHandle = CreateThread(
    NULL, // default security attributes
    0, // default stack size
    Summation, // thread function
    &Param, // parameter to thread function
    0, // default creation flags
    &ThreadId); // returns the thread identifier

if (ThreadHandle != NULL) {
    // now wait for the thread to finish
    WaitForSingleObject(ThreadHandle, INFINITE);

    // close the thread handle
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}
```

Figure 4.10 Multithreaded C program using the Win32 API.

Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```


Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                // create the object to be shared
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of " + upper + " is " + sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>"); }
}
```

Figure 4.11 Java program for the summation of a non-negative integer.